

```

import numpy as np
import gym
from collections import deque
import random

# Ornstein-Uhlenbeck Process
# Taken from
# https://github.com/vitchyr/rlkit/blob/master/rlkit/exploration\_strategies/ou\_strategy.py
class OUNoise(object):
    def __init__(self, action_space, mu=0.0, theta=0.15,
max_sigma=0.3, min_sigma=0.3, decay_period=100000):
        self.mu = mu
        self.theta = theta
        self.sigma = max_sigma
        self.max_sigma = max_sigma
        self.min_sigma = min_sigma
        self.decay_period = decay_period
        self.action_dim = action_space.shape[0]
        self.low = action_space.low
        self.high = action_space.high
        self.reset()

    def reset(self):
        self.state = np.ones(self.action_dim) * self.mu

    def evolve_state(self):
        x = self.state
        dx = self.theta * (self.mu - x) + self.sigma *
np.random.randn(self.action_dim)
        self.state = x + dx
        return self.state

    def get_action(self, action, t=0):
        ou_state = self.evolve_state()
        self.sigma = self.max_sigma - (self.max_sigma -
self.min_sigma) * min(1.0, t / self.decay_period)
        return np.clip(action + ou_state, self.low, self.high)

# https://github.com/openai/gym/blob/master/gym/core.py
class NormalizedEnv(gym.ActionWrapper):
    """ Wrap action """

    def action(self, action):
        act_k = (self.action_space.high - self.action_space.low) / 2.
        act_b = (self.action_space.high + self.action_space.low) / 2.
        return act_k * action + act_b

```

```

class Memory:
    def __init__(self, max_size):
        self.max_size = max_size
        self.buffer = deque(maxlen=max_size)

    def push(self, state, action, reward, next_state, done):
        experience = (state, action, np.array([reward]), next_state,
done)
        self.buffer.append(experience)

    def sample(self, batch_size):
        state_batch = []
        action_batch = []
        reward_batch = []
        next_state_batch = []
        done_batch = []

        batch = random.sample(self.buffer, batch_size)

        for experience in batch:
            state, action, reward, next_state, done = experience
            state_batch.append(state)
            action_batch.append(action)
            reward_batch.append(reward)
            next_state_batch.append(next_state)
            done_batch.append(done)

        return state_batch, action_batch, reward_batch,
next_state_batch, done_batch

    def __len__(self):
        return len(self.buffer)

```

DDPG uses four neural networks: a Q network, a deterministic policy network, a target Q network, and a target policy network.

## Parameters:

$\theta^Q$  : Q network

$\theta^\mu$  : Deterministic policy function

$\theta^{Q'}$  : target Q network

$\theta^{\mu'}$  : target policy network

The Q network and policy network is very much like simple Advantage Actor-Critic, but in DDPG, the Actor directly maps states to actions instead of outputting the probability distribution across a discrete action space.

The target networks are time-delayed copies of their original networks that slowly track the learned networks. Using these target value networks greatly improve stability in learning.

Let's create these networks.

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class Critic(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(Critic, self).__init__()
        self.linear1 = nn.Linear(input_size, hidden_size)
        self.linear2 = nn.Linear(hidden_size, hidden_size)
        self.linear3 = nn.Linear(hidden_size, output_size)

    def forward(self, state, action):
        """
        Params state and actions are torch tensors
        """
        x = torch.cat([state, action], 1)
        x = F.relu(self.linear1(x))
        x = F.relu(self.linear2(x))
        x = self.linear3(x)

        return x

class Actor(nn.Module):
```

```

def __init__(self, input_size, hidden_size, output_size,
learning_rate = 3e-4):
    super(Actor, self).__init__()
    self.linear1 = nn.Linear(input_size, hidden_size)
    self.linear2 = nn.Linear(hidden_size, hidden_size)
    self.linear3 = nn.Linear(hidden_size, output_size)

def forward(self, state):
    """
    Param state is a torch tensor
    """
    x = F.relu(self.linear1(state))
    x = F.relu(self.linear2(x))
    x = torch.tanh(self.linear3(x))

    return x

```

Now, let's create the DDPG agent. The agent class has two main functions: "get\_action" and "update":

- **get\_action()**: This function runs a forward pass through the actor network to select a deterministic action. In the DDPG paper, the authors use Ornstein-Uhlenbeck Process to add noise to the action output (Uhlenbeck & Ornstein, 1930), thereby resulting in exploration in the environment. Class OUNoise (in cell 1) implements this.

$$\mu'(s_t) = \mu(s_t | \theta_t^\mu) + \mathcal{N}$$

- **update()**: This function is used for updating the actor and critic networks, and forms the core of the DDPG algorithm. The replay buffer is first sampled to get a batch of experiences of the form **<states, actions, rewards, next\_states>**.

The value network is updated using the Bellman equation, similar to Q-learning. However, in DDPG, the next-state Q values are calculated with the target value network and target policy network. Then, we minimize the mean-squared loss between the target Q value and the predicted Q value:

$$y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1} | \theta^{\mu'}) | \theta^{Q'})$$

$$Loss = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i | \theta^Q))^2$$

For the policy function, our objective is to maximize the expected return. To calculate the policy gradient, we take the derivative of the objective function with respect to the policy parameter. For this, we use the chain rule.

$$\nabla_{\theta^\mu} J(\theta) \approx \frac{1}{N} \sum_i [\nabla_a Q(s, a | \theta^Q) |_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s | \theta^\mu) |_{s=s_i}]$$

We make a copy of the target network parameters and have them slowly track those of the learned networks via “soft updates,” as illustrated below:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

$$\text{where } \tau \ll 1$$

```
import torch
import torch.optim as optim
import torch.nn as nn

class DDPGagent:
    def __init__(self, env, hidden_size=256, actor_learning_rate=1e-4,
critic_learning_rate=1e-3, gamma=0.99, tau=1e-2,
max_memory_size=50000):
        # Params
        self.num_states = env.observation_space.shape[0]
        self.num_actions = env.action_space.shape[0]
        self.gamma = gamma
        self.tau = tau

        # Networks
        self.actor = Actor(self.num_states, hidden_size,
self.num_actions)
        self.actor_target = Actor(self.num_states, hidden_size,
self.num_actions)
        self.critic = Critic(self.num_states + self.num_actions,
hidden_size, self.num_actions)
        self.critic_target = Critic(self.num_states +
self.num_actions, hidden_size, self.num_actions)
```

```

        for target_param, param in zip(self.actor_target.parameters(),
self.actor.parameters()):
            target_param.data.copy_(param.data)

        for target_param, param in
zip(self.critic_target.parameters(), self.critic.parameters()):
            target_param.data.copy_(param.data)

    # Training
    self.memory = Memory(max_memory_size)
    self.critic_criterion = nn.MSELoss()
    self.actor_optimizer = optim.Adam(self.actor.parameters(),
lr=actor_learning_rate)
    self.critic_optimizer = optim.Adam(self.critic.parameters(),
lr=critic_learning_rate)

    def get_action(self, state):
        state = torch.FloatTensor(state).unsqueeze(0)
        action = self.actor.forward(state)
        action = action.detach().numpy()[0,0]
        return action

    def update(self, batch_size):
        states, actions, rewards, next_states, _ =
self.memory.sample(batch_size)
        states = torch.FloatTensor(states)
        actions = torch.FloatTensor(actions)
        rewards = torch.FloatTensor(rewards)
        next_states = torch.FloatTensor(next_states)

        # Implement critic loss and update critic
        self.critic_optimizer.zero_grad()
        y = rewards + self.gamma * self.critic_target.forward(states,
self.actor_target.forward(states).detach())
        Q = self.critic.forward(states, actions)
        loss = self.critic_criterion(y, Q)
        loss.backward()
        self.critic_optimizer.step()

        # Implement actor loss and update actor
        self.actor_optimizer.zero_grad()
        policy_loss = -self.critic.forward(states,
self.actor.forward(states))
        policy_loss = policy_loss.mean()
        policy_loss.backward()
        self.actor_optimizer.step()

    # update target networks

```

```

        with torch.no_grad():
            for target_param, param in
zip(self.actor_target.parameters(), self.actor.parameters()):
                target_param.data.copy_(self.tau * param.data + (1.0 -
self.tau) * target_param.data)
            for target_param, param in
zip(self.critic_target.parameters(), self.critic.parameters()):
                target_param.data.copy_(self.tau * param.data + (1.0 -
self.tau) * target_param.data)

```

Putting it all together: DDPG in action.

The main function below runs 100 episodes of DDPG on the "Pendulum-v0" environment of OpenAI gym. This is the inverted pendulum swingup problem, a classic problem in the control literature. In this version of the problem, the pendulum starts in a random position, and the goal is to swing it up so it stays upright.

Each episode is for a maximum of 200 timesteps. At each step, the agent chooses an action, moves to the next state and updates its parameters according to the DDPG algorithm, repeating this process till the end of the episode.

The DDPG algorithm is as follows:

---

**Algorithm 1** DDPG algorithm

---

Randomly initialize critic network  $Q(s, a|\theta^Q)$  and actor  $\mu(s|\theta^\mu)$  with weights  $\theta^Q$  and  $\theta^\mu$ .  
Initialize target network  $Q'$  and  $\mu'$  with weights  $\theta^{Q'} \leftarrow \theta^Q$ ,  $\theta^{\mu'} \leftarrow \theta^\mu$   
Initialize replay buffer  $R$   
**for** episode = 1,  $M$  **do**  
    Initialize a random process  $\mathcal{N}$  for action exploration  
    Receive initial observation state  $s_1$   
    **for**  $t = 1, T$  **do**  
        Select action  $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$  according to the current policy and exploration noise  
        Execute action  $a_t$  and observe reward  $r_t$  and observe new state  $s_{t+1}$   
        Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $R$   
        Sample a random minibatch of  $N$  transitions  $(s_i, a_i, r_i, s_{i+1})$  from  $R$   
        Set  $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$   
        Update critic by minimizing the loss:  $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$   
        Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

**end for**  
**end for**

---

```

import sys
import gym
import numpy as np

```

```

import pandas as pd
import matplotlib.pyplot as plt

# For more info on the Pendulum environment, check out
https://www.gymnasium.dev/environments/classic\_control/pendulum/
env = NormalizedEnv(gym.make("Pendulum-v1"))

agent = DDPGAgent(env)
noise = OUNoise(env.action_space)
batch_size = 128
rewards = []
avg_rewards = []

for episode in range(100):
    state = env.reset()
    noise.reset()
    episode_reward = 0

    for step in range(200):
        action = agent.get_action(state)
        #Add noise to action

        action = noise.get_action(action, step)
        new_state, reward, done, _ = env.step(action)
        agent.memory.push(state, action, reward, new_state, done)

        if len(agent.memory) > batch_size:
            agent.update(batch_size)

        state = new_state
        episode_reward += reward

    if done:
        sys.stdout.write("episode: {}, reward: {}, average
_reward: {} \n".format(episode, np.round(episode_reward, decimals=2),
np.mean(rewards[-10:])))
        break

    rewards.append(episode_reward)
    avg_rewards.append(np.mean(rewards[-10:]))

plt.plot(rewards)
plt.plot(avg_rewards)
plt.plot()
plt.xlabel('Episode')
plt.ylabel('Reward')
plt.show()

```

/usr/local/lib/python3.9/dist-packages/gym/core.py:317:  
DeprecationWarning: WARN: Initializing wrapper in old step API which



returns one bool instead of two. It is recommended to set  
`new\_step\_api=True` to use new step API. This will be the default  
behaviour in future.

```
deprecation(  
/usr/local/lib/python3.9/dist-packages/gym/wrappers/step_api_compatibi  
lity.py:39: DeprecationWarning: WARN: Initializing environment in old  
step API which returns one bool instead of two. It is recommended to  
set `new_step_api=True` to use new step API. This will be the default  
behaviour in future.
```

```
deprecation(  
/usr/local/lib/python3.9/dist-packages/numpy/core/fromnumeric.py:3474:  
RuntimeWarning: Mean of empty slice.
```

```
return _methods._mean(a, axis=axis, dtype=dtype,  
/usr/local/lib/python3.9/dist-packages/numpy/core/_methods.py:189:  
RuntimeWarning: invalid value encountered in double_scalars  
ret = ret.dtype.type(ret / rcount)
```

```
episode: 0, reward: -1402.47, average _reward: nan  
episode: 1, reward: -1395.09, average _reward: -1402.4658915334019  
episode: 2, reward: -1727.51, average _reward: -1398.7804329776304  
episode: 3, reward: -1441.61, average _reward: -1508.3581090076798  
episode: 4, reward: -1543.35, average _reward: -1491.6700328552874  
episode: 5, reward: -1256.06, average _reward: -1502.005737924385  
episode: 6, reward: -1557.2, average _reward: -1461.014589812629  
episode: 7, reward: -1593.07, average _reward: -1474.7547231347285  
episode: 8, reward: -1538.75, average _reward: -1489.5439029584786  
episode: 9, reward: -1308.45, average _reward: -1495.01114188601  
episode: 10, reward: -1567.27, average _reward: -1476.3549182694662  
episode: 11, reward: -1410.96, average _reward: -1492.8352379433177  
episode: 12, reward: -1404.76, average _reward: -1494.4217679105834  
episode: 13, reward: -1475.55, average _reward: -1462.1465197359962  
episode: 14, reward: -1356.35, average _reward: -1465.540560362231  
episode: 15, reward: -1380.27, average _reward: -1446.8407785223942  
episode: 16, reward: -1470.51, average _reward: -1459.2615009091967  
episode: 17, reward: -1522.62, average _reward: -1450.5930749253298  
episode: 18, reward: -1387.59, average _reward: -1443.548709983089  
episode: 19, reward: -1263.18, average _reward: -1428.433049155354  
episode: 20, reward: -1387.72, average _reward: -1423.906656520531  
episode: 21, reward: -1305.01, average _reward: -1405.9518878826525  
episode: 22, reward: -1000.29, average _reward: -1395.3567623897322  
episode: 23, reward: -1159.95, average _reward: -1354.9101213649794  
episode: 24, reward: -1465.45, average _reward: -1323.3506531170342  
episode: 25, reward: -1155.1, average _reward: -1334.2605453256697  
episode: 26, reward: -1187.53, average _reward: -1311.7439829404107  
episode: 27, reward: -1362.4, average _reward: -1283.4455385137169  
episode: 28, reward: -1334.57, average _reward: -1267.423073804179  
episode: 29, reward: -1239.95, average _reward: -1262.121088679628  
episode: 30, reward: -1237.3, average _reward: -1259.7980011625828  
episode: 31, reward: -1427.09, average _reward: -1244.756040747589  
episode: 32, reward: -1381.09, average _reward: -1256.964093349837
```

episode: 33, reward: -1387.82, average \_reward: -1295.0431779729956  
episode: 34, reward: -1291.11, average \_reward: -1317.8302865021265  
episode: 35, reward: -1188.46, average \_reward: -1300.3965294941895  
episode: 36, reward: -1045.18, average \_reward: -1303.7319897171697  
episode: 37, reward: -1189.55, average \_reward: -1289.4972926913563  
episode: 38, reward: -1451.08, average \_reward: -1272.2121308049334  
episode: 39, reward: -1213.42, average \_reward: -1283.8624108792985  
episode: 40, reward: -1437.29, average \_reward: -1281.208642125628  
episode: 41, reward: -1400.69, average \_reward: -1301.2069969611844  
episode: 42, reward: -1213.43, average \_reward: -1298.5673840508387  
episode: 43, reward: -1404.42, average \_reward: -1281.8019115894078  
episode: 44, reward: -1419.48, average \_reward: -1283.4620612114861  
episode: 45, reward: -1233.36, average \_reward: -1296.2989561982718  
episode: 46, reward: -1270.65, average \_reward: -1300.789819911657  
episode: 47, reward: -1199.6, average \_reward: -1323.3369189012717  
episode: 48, reward: -1159.46, average \_reward: -1324.3423466949273  
episode: 49, reward: -975.94, average \_reward: -1295.1809773318835  
episode: 50, reward: -1315.01, average \_reward: -1271.432993044808  
episode: 51, reward: -1273.25, average \_reward: -1259.2050369292438  
episode: 52, reward: -1061.19, average \_reward: -1246.461049140358  
episode: 53, reward: -1217.33, average \_reward: -1231.2367259256089  
episode: 54, reward: -1390.65, average \_reward: -1212.5270508378362  
episode: 55, reward: -1318.47, average \_reward: -1209.6440997920377  
episode: 56, reward: -1329.55, average \_reward: -1218.1550078874575  
episode: 57, reward: -1270.72, average \_reward: -1224.0451091919049  
episode: 58, reward: -1463.07, average \_reward: -1231.1569382211271  
episode: 59, reward: -1274.7, average \_reward: -1261.5179627913833  
episode: 60, reward: -1355.71, average \_reward: -1291.3947416029573  
episode: 61, reward: -1431.16, average \_reward: -1295.4653690953005  
episode: 62, reward: -928.18, average \_reward: -1311.2555935486691  
episode: 63, reward: -1275.4, average \_reward: -1297.954784031932  
episode: 64, reward: -915.49, average \_reward: -1303.7618379893488  
episode: 65, reward: -1413.34, average \_reward: -1256.2461817151782  
episode: 66, reward: -1371.74, average \_reward: -1265.7327449316567  
episode: 67, reward: -1395.45, average \_reward: -1269.9511909172227  
episode: 68, reward: -1204.49, average \_reward: -1282.4236678726763  
episode: 69, reward: -1413.36, average \_reward: -1256.5654100033294  
episode: 70, reward: -1343.45, average \_reward: -1270.43093486687  
episode: 71, reward: -1291.35, average \_reward: -1269.2046400473473  
episode: 72, reward: -1409.47, average \_reward: -1255.2236614176068  
episode: 73, reward: -1153.02, average \_reward: -1303.3528936099221  
episode: 74, reward: -1241.89, average \_reward: -1291.1146402519355  
episode: 75, reward: -1389.24, average \_reward: -1323.75460204643  
episode: 76, reward: -1307.99, average \_reward: -1321.3452203841516  
episode: 77, reward: -1262.11, average \_reward: -1314.970273786325  
episode: 78, reward: -1163.74, average \_reward: -1301.6370330687398  
episode: 79, reward: -1210.1, average \_reward: -1297.5618793822541  
episode: 80, reward: -1254.75, average \_reward: -1277.2358084927241  
episode: 81, reward: -1367.09, average \_reward: -1268.3656534736106  
episode: 82, reward: -1278.97, average \_reward: -1275.93994060668

episode: 83, reward: -1431.52, average\_reward: -1262.8895849595451  
episode: 84, reward: -1377.3, average\_reward: -1290.73963987675  
episode: 85, reward: -1240.16, average\_reward: -1304.2801330915845  
episode: 86, reward: -1178.59, average\_reward: -1289.3712245549063  
episode: 87, reward: -1094.93, average\_reward: -1276.4320306751958  
episode: 88, reward: -1489.05, average\_reward: -1259.7132660678985  
episode: 89, reward: -1394.83, average\_reward: -1292.2448904274777  
episode: 90, reward: -1256.37, average\_reward: -1310.7179932806673  
episode: 91, reward: -1408.08, average\_reward: -1310.8799556781396  
episode: 92, reward: -1258.31, average\_reward: -1314.9791657968058  
episode: 93, reward: -1178.55, average\_reward: -1312.9130992630771  
episode: 94, reward: -1163.27, average\_reward: -1287.6160523416108  
episode: 95, reward: -1214.84, average\_reward: -1266.2129203235613  
episode: 96, reward: -1193.66, average\_reward: -1263.6808822552334  
episode: 97, reward: -1316.45, average\_reward: -1265.1870229422207  
episode: 98, reward: -1223.69, average\_reward: -1287.3395048303385  
episode: 99, reward: -1267.18, average\_reward: -1260.803433692495

